

Boosting

Seminar: Foundations of Data Science

Christian Peters, enrolment no.: 213996
Faculty of Statistics
TU Dortmund

January 11, 2021
winter term 2020/21

Abstract

Boosting is an algorithmic paradigm that aims to create an efficient strong learner by combining multiple weak learners. This paper introduces the concept of weak learnability and explains the AdaBoost algorithm, the first practical implementation of boosting.

1 Introduction

Training machine learning models in practice is not always as simple as it might seem from a theoretical standpoint. In [1, chapter 2], the empirical risk minimization (ERM) rule was introduced as the learning algorithm of choice. However, this theoretical principle of choosing a hypothesis $h \in \mathcal{H}$ such that $L_S(h) = \min_{h \in \mathcal{H}} L_S(h)$, where $L_S(h)$ is the error of h on the training sequence S , can be impossible to use in practical applications due to the sometimes enormous computational complexity of searching through interesting hypothesis classes \mathcal{H} .

This problem leads to the question if it is possible to arrive at a strong learning algorithm in a way that doesn't require the computational cost of searching through complex hypothesis classes. Is it perhaps possible to create a strong learner by finding a way to combine "weak" learners that are potentially easier to compute? The algorithmic paradigm of boosting deals with exactly this question, resulting in the widely used AdaBoost algorithm that shows how "weak" learners can be combined in order to obtain a strong learning algorithm.

23 The goal of this article is to first lay down the foundations of boosting
 24 by explaining the concept of weak learnability which will be used to arrive
 25 at the AdaBoost algorithm followed by a discussion of its implications as
 26 well as a practical example of how it can be used in the domain of image
 27 classification.

28 2 Weak Learnability

29 Assuming that the realizable assumption [1, chapter 2] holds, the generaliza-
 30 tion error $L_{(\mathcal{D},f)}(h)$ of a PAC learner with respect to a distribution \mathcal{D} and
 31 a labeling function f can by the definition of PAC learning [1, chapter 2]
 32 be reduced to an arbitrarily small number (with confidence of $1 - \delta$) by in-
 33 creasing the sample size of the training sequence S . This, however, can be
 34 computationally infeasible in practical applications.

35 The concept of weak learnability aims to relax the requirement that the
 36 generalization error of a hypothesis must become arbitrarily small the more
 37 the sample size is increased. For weak learning, it is sufficient that the
 38 learning algorithm yields a hypothesis h , that performs only slightly better
 39 than a random guess. One can think of weak learning as applying a simple
 40 rule which isn't fully capable of modeling the data generating process, but
 41 can still learn a little bit about the underlying problem so that it performs
 42 better than random.

43 Formally, the definition of a weak learner differs only slightly from the
 44 definition of PAC learning. In the situation of a two-class classification prob-
 45 lem, the definition of a weak learner, can be given as follows:

46 **Definition 1.** (*γ -weak-learner*) An Algorithm A is a γ -weak-learner for a
 47 hypothesis class \mathcal{H} if for every $\delta \in (0, 1)$ there exists a threshold $m_{\mathcal{H}}(\delta) \in \mathbb{N}$
 48 such that for every distribution \mathcal{D} over the instance space \mathcal{X} and for every
 49 labeling function $f : \mathcal{X} \rightarrow \{\pm 1\}$ if running A on $m \geq m_{\mathcal{H}}$ training examples,
 50 it will yield a hypothesis h such that with probability of at least $1 - \delta$ the
 51 generalization error $L_{(\mathcal{D},f)}(h)$ is at most $\frac{1}{2} - \gamma$, provided that the realizable
 52 assumption holds.

53 The parameter γ in Definition 1 tells us how well we can expect the weak
 54 learner to perform. For example if a learning algorithm is a 1%-weak-learner
 55 for a class \mathcal{H} , then the generalization error can at most be 49% provided that
 56 first, the learner didn't fail (which can happen with probability δ of drawing
 57 a bad sample from \mathcal{D}) and second, the learner was run on at least $m_{\mathcal{H}}(\delta)$
 58 training examples.

59 It still remains the question of how to obtain a weak learning algorithm
60 for a class \mathcal{H} . We already saw that applying the ERM rule to complex hy-
61 pothesis classes can be computationally hard. But what if we choose a simple
62 hypothesis class B instead, where the ERM rule can be applied efficiently?
63 It follows directly from Definition 1 that this can only work if the new algo-
64 rithm ERM_B has an error of at most $\frac{1}{2} - \gamma$ for every sample that was labeled
65 by a hypothesis from \mathcal{H} . In this case, applying the ERM rule with respect
66 to the simpler class B would yield a weak learner for \mathcal{H} .

67 2.1 An Efficient ERM Algorithm for Decision Stumps

68 One such hypothesis class, where the ERM algorithm can be implemented
69 efficiently, is the class of decision stumps. On the instance space $\mathcal{X} = \mathbb{R}^d$,
70 this class is given as follows:

$$71 \quad \mathcal{H}_{DS} = \{\mathbf{x} \mapsto \text{sign}(\theta - x_i) \cdot b : \theta \in \mathbb{R}, i \in [d], b \in \{\pm 1\}\}$$

72 The idea behind decision stumps is to divide the instance space \mathcal{X} along
73 one of its dimensions $i \in [d]$ at a threshold θ into two partitions, such that
74 all the instances in one partition are labeled +1 and all the instances in the
75 other partition are labeled -1.

76 Fixing $b = 1$ without the loss of generality, an ERM algorithm for \mathcal{H}_{DS}
77 has to find the optimal splitting dimension $i \in [d]$ as well as the optimal
78 splitting threshold θ such that the training error $L_S(h)$ on a training sequence
79 $S = ((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m))$ is minimized.

80 An extension of the empirical risk $L_S(h)$ that will be useful when intro-
81 ducing the AdaBoost algorithm later in section 3, is the weighted empirical
82 risk, which is given as

$$83 \quad L_{\mathbf{D}}(h) = \sum_{i=1}^m D_i \mathbb{1}_{[h(\mathbf{x}_i) \neq y_i]}.$$

84 The vector $\mathbf{D} \in \mathbb{R}^m$ is used to give a weight to each training example, where
85 each weight D_i is nonnegative and $\sum_{i=1}^m D_i = 1$. For the special case that
86 $D_i = \frac{1}{m}$, the weighted empirical risk is equal to the unweighted empirical
87 risk.

88 In order to find a decision stump $h \in \mathcal{H}_{DS}$ that minimizes the weighted
89 empirical risk $L_{\mathbf{D}}(h)$, the $\text{ERM}_{\mathcal{H}_{DS}}$ algorithm has to solve the following op-
90 timization problem that minimizes the weighted sum of misclassifications:

$$91 \quad \min_{j \in [d]} \min_{\theta \in \mathbb{R}} \left(\sum_{i: y_i = 1}^m D_i \mathbb{1}_{[x_{i,j} > \theta]} + \sum_{i: y_i = -1}^m D_i \mathbb{1}_{[x_{i,j} \leq \theta]} \right)$$

92 Taking a closer look at this problem it becomes clear, that it is not necessary
 93 to try every $\theta \in \mathbb{R}$ during the optimization. In fact, if we first sort the
 94 examples for a fixed dimension $j \in [d]$ so that $x_{1,j} \leq x_{2,j} \leq \dots \leq x_{m,j}$, we
 95 can see that we only have to consider a single splitting threshold in between
 96 two examples, which leaves us with $m + 1$ values for θ that the $\text{ERM}_{\mathcal{H}_{DS}}$ has
 97 to consider for a fixed dimension $j \in [d]$. Since the sum of misclassifications
 98 can be computed in $\mathcal{O}(m)$ by passing through the data once, the algorithm
 99 can find optimal values for j and θ in $\mathcal{O}(dm^2)$ by simple enumeration.

100 This time complexity can be reduced even more by avoiding to recalculate
 101 the sum of misclassifications for every new value of θ . It can be shown [1] that
 102 it is possible to update the sum for a new value of θ in constant time, requiring
 103 only a single pass through the data for a fixed dimension j . This leaves us
 104 with a time complexity of $\mathcal{O}(dm)$ of solving the optimization problem after
 105 the sorting step is applied as preprocessing.

106 **3 AdaBoost**

107 In the previous section it was shown how an efficient weak learner can be
 108 constructed by applying the ERM rule to the class of decision stumps. This
 109 section presents the AdaBoost (Adaptive Boosting) algorithm, a procedure
 110 that shows how weak learners such as decision stumps can be used efficiently
 111 to find a hypothesis with an arbitrarily low empirical error $L_S(h)$ on a training
 112 sequence S .

113 The goal of AdaBoost is to invoke the weak learner multiple times on
 114 the training data and then to combine the resulting hypotheses similar to a
 115 weighted majority vote. Let T be the number of times that AdaBoost invokes
 116 the weak learner on the training data. Then the resulting output hypothesis
 117 of AdaBoost has the following form:

$$118 \quad h(x) = \text{sign} \left(\sum_{t=1}^T w_t h_t(x) \right)$$

119 Here, $h_t(x)$ is the output hypothesis of the weak learner in iteration t and w_t
 120 is the corresponding weight that AdaBoost assigns to this hypothesis.

121 In the first iteration $t = 1$, the AdaBoost algorithm assigns an equal
 122 weight $D_i^{(1)} = \frac{1}{m}$ to each example in the training sequence S and then invokes
 123 the weak learner on the weighted training sequence. The error of the resulting
 124 hypothesis is computed according to

$$125 \quad \epsilon_t = \sum_{i=1}^m D_i^{(t)} \mathbb{1}_{[h_t(\mathbf{x}_i) \neq y_i]}$$

126 and is at most $\frac{1}{2} - \gamma$.

127 The weight w_t that is assigned to a hypothesis in boosting round t is
128 computed as follows:

$$129 \quad w_t = \frac{1}{2} \log \left(\frac{1}{\epsilon_t} - 1 \right)$$

130 As we can see, the smaller the error ϵ_t of the hypothesis h_t is, the bigger the
131 weight w_t will be.

132 At the end of each iteration, the weights $D_i^{(t)}$ of each training example
133 are updated according to

$$134 \quad D_i^{(t+1)} = \frac{D_i^{(t)} \exp(-w_t y_i h_t(\mathbf{x}_i))}{\sum_{j=1}^m D_j^{(t)} \exp(-w_t y_j h_t(\mathbf{x}_j))}$$

135 This update assigns a higher weight to those examples, that weren't correctly
136 classified by h_t .

137 In short, the Adaboost algorithm performs the following steps in each of
138 the T iterations:

- 139 1. Invoke the weak learner on the training sequence S weighted by $\mathbf{D}^{(t)}$
- 140 2. Assign a weight w_t to the output hypothesis h_t of the weak learner.
141 Hypotheses with a smaller training error ϵ_t will get a higher weight.
- 142 3. Compute a new weight vector $\mathbf{D}^{(t+1)}$ that gives a higher weight to
143 incorrectly classified examples.

144 The computational complexity of AdaBoost essentially consists of invoking
145 the weak learning algorithm T times on the training data. Thus, if the
146 weak learner can be implemented efficiently (as it is the case for decision
147 stumps), AdaBoost is also efficient.

148 On top of that, it can be shown [1], that the training error $L_S(h)$ of the
149 AdaBoost hypothesis h decreases exponentially in the number of boosting
150 rounds T :

$$151 \quad L_S(h) = \frac{1}{m} \sum_{i=1}^m \mathbb{1}_{[h(\mathbf{x}_i) \neq y_i]} \leq e^{-2\gamma^2 T}$$

152 Here, γ describes the γ -weak learner as defined in Definition 1. This means,
153 that AdaBoost will achieve an arbitrarily low training error and is still an
154 efficient algorithm.

155 In practical applications however, it is more important to also achieve
156 a good out of sample error. The next section will show that the true risk
157 $L_{\mathcal{D}}(h)$ of the AdaBoost hypothesis will also be small by taking a look at
158 the hypothesis class that resembles all the possible output hypotheses of
159 AdaBoost.

160 3.1 AdaBoost Out-of-Sample Performance

161 The hypothesis class of AdaBoost is parameterized by the hypothesis class
 162 B of the weak learner it uses as well as by the number of boosting rounds T .
 163 Formally, it is given as follows:

$$164 \quad L(B, T) = \left\{ x \mapsto \text{sign} \left(\sum_{t=1}^T w_t h_t(x) \right) : w \in \mathbb{R}^T, h_t \in B \right\}$$

165 The fundamental theorem of statistical learning [1, chapter 6] states
 166 that a hypothesis class is PAC-learnable if its VC dimension is finite. This
 167 means that if the VC dimension of $L(B, T)$ is finite, there exists a threshold
 168 $m_{\mathcal{H}}(\epsilon, \delta) \in \mathbb{N}$ for every $\epsilon, \delta \in (0, 1)$, such that $L_{\mathcal{D}}(L(B, T)) \leq \epsilon$ (with con-
 169 fidence $1 - \delta$) for every distribution \mathcal{D} when training AdaBoost on at least
 170 $m_{\mathcal{H}}(\epsilon, \delta)$ training examples. This tells us that even the Out-of-Sample error
 171 of AdaBoost can be reduced arbitrarily (if the realizable assumption holds),
 172 by increasing the amount of training examples.

173 It can be shown that an upper bound of the VC dimension of $L(B, T)$
 174 is linear in the VC dimension of the hypothesis class B of the weak learner
 175 and also linear in T . This means that if the VC dimension of B is finite, so
 176 is the VC dimension of $L(B, T)$.

177 In the case of B being the hypothesis class of all decision stumps, the
 178 VC dimension of B is 2, so it is clearly finite. It follows from the funda-
 179 mental theorem of statistical learning, that when using decision stumps as
 180 the hypothesis class for the weak learner, AdaBoost is a PAC learner for the
 181 class $L(B, T)$. Furthermore, if the weak learner for B can be implemented
 182 efficiently (as we have shown for decision stumps), AdaBoost is also efficient.

183 4 Conclusion

184 Even after centuries of research in the field of data science, there is nothing
 185 more versatile than the useful theorem of chapter 3. It is used everywhere
 186 and has led to the greatest and most intriguing results, cf. [2]. By the way,
 187 the book for the seminar [1] is a great reference and should be cited. Further
 188 literature can be found in the respective *Bibliographic Remarks* sections and
 189 of course you are welcome to search and add your own references.

190 **Note:** BibTeX entries can often be found in the DBLP collection. Google
 191 Scholar also offers BibTeX entries, which can be copied into the .bib file and
 192 may need some minor adjustments.

References

- [1] S. Shalev-Shwartz and S. Ben-David. *Understanding Machine Learning - From Theory to Algorithms*. Cambridge University Press, 2014.
- [2] J. Someone and J. Someoneelse. Useful theorems, 2003.