

Boosting

Seminar: Foundations of Data Science

Christian Peters, enrolment no.: 213996
Faculty of Statistics
TU Dortmund

December 30, 2020
winter term 2020/21

Abstract

A short summary of about two to three sentences that briefly and
concisely outline the content ...

1 Introduction

Training machine learning models in practice is not always as simple as it might seem from a theoretical standpoint. In [1, chapter 2], the empirical risk minimization (ERM) rule was introduced as the learning algorithm of choice. However, this theoretical principle of choosing a hypothesis $h \in \mathcal{H}$ such that $L_S(h) = \min_{h \in \mathcal{H}} L_S(h)$, where $L_S(h)$ is the error of h on the training sequence S , can be impossible to use in practical applications due to the sometimes enormous computational complexity of searching through interesting hypothesis classes \mathcal{H} .

This problem leads to the question if it is possible to arrive at a strong learning algorithm in a way that doesn't require the computational cost of searching through complex hypothesis classes. Is it perhaps possible to create a strong learner by finding a way to combine "weak" learners that are potentially easier to compute? The algorithmic paradigm of boosting deals with exactly this question, resulting in the widely used AdaBoost algorithm that shows how "weak" learners can be combined in order to obtain a strong learning algorithm.

The goal of this article is to first lay down the foundations of boosting by explaining the concept of weak learnability which will be used to arrive

23 at the AdaBoost algorithm followed by a discussion of its implications as
 24 well as a practical example of how it can be used in the domain of image
 25 classification.

26 2 Weak Learnability

27 Assuming that the realizable assumption [1, chapter 2] holds, the generaliza-
 28 tion error $L_{(\mathcal{D},f)}(h)$ of a PAC learner with respect to a distribution \mathcal{D} and
 29 a labeling function f can by the definition of PAC learning [1, chapter 2]
 30 be reduced to an arbitrarily small number (with confidence of $1 - \delta$) by in-
 31 creasing the sample size of the training sequence S . This, however, can be
 32 computationally infeasible in practical applications.

33 The concept of weak learnability aims to relax the requirement that the
 34 generalization error of a hypothesis must become arbitrarily small the more
 35 the sample size is increased. For weak learning, it is sufficient that the
 36 learning algorithm yields a hypothesis h , that performs only slightly better
 37 than a random guess. One can think of weak learning as applying a simple
 38 rule which isn't fully capable of modeling the data generating process, but
 39 can still learn a little bit about the underlying problem so that it performs
 40 better than random.

41 Formally, the definition of a weak learner differs only slightly from the
 42 definition of PAC learning. In the situation of a two-class classification prob-
 43 lem, the definition of a weak learner, can be given as follows:

44 **Definition 1.** (*γ -weak-learner*) An Algorithm A is a γ -weak-learner for a
 45 hypothesis class \mathcal{H} if for every $\delta \in (0, 1)$ there exists a threshold $m_{\mathcal{H}}(\delta) \in \mathbb{N}$
 46 such that for every distribution \mathcal{D} over the instance space \mathcal{X} and for every
 47 labeling function $f : \mathcal{X} \rightarrow \{\pm 1\}$ if running A on $m \geq m_{\mathcal{H}}$ training examples,
 48 it will yield a hypothesis h such that with probability of at least $1 - \delta$ the
 49 generalization error $L_{(\mathcal{D},f)}(h)$ is at most $\frac{1}{2} - \gamma$, provided that the realizable
 50 assumption holds.

51 The parameter γ in Definition 1 tells us how well we can expect the weak
 52 learner to perform. For example if a learning algorithm is a 1%-weak-learner
 53 for a class \mathcal{H} , then the generalization error can at most be 49% provided that
 54 first, the learner didn't fail (which can happen with probability δ of drawing
 55 a bad sample from \mathcal{D}) and second, the learner was run on at least $m_{\mathcal{H}}(\delta)$
 56 training examples.

57 It still remains the question of how to obtain a weak learning algorithm.
 58 We already saw that applying the ERM rule to complex hypothesis classes

can be computationally hard. But what if we choose a simple hypothesis class B instead, where the ERM rule can be applied efficiently?

Theorem 1. *Let B and \mathcal{H} be hypothesis classes and ERM_B a learning algorithm that applies the ERM rule to B . If for every distribution \mathcal{D} , for every $h \in \mathcal{H}$ and for every sample $S = ((x_1, h(x_1)), \dots, (x_m, h(x_m)))$ from \mathcal{D} labeled by h it holds that ERM_B returns a hypothesis $b \in B$ with $L_S(b) \leq \frac{1}{2} - \gamma$, then ERM_B is a γ -weak-learner for \mathcal{H} .*

Proof. TODO. □

3 AdaBoost

Here we work with the above definitions and notations and derive important results such as the following Theorem on the least-squares solution

Theorem 2. *(useful theorem) Let $X \in \mathbb{R}^{n \times d}$, $Y \in \mathbb{R}^n$. Further define $\beta^* = \operatorname{argmin}_{\beta \in \mathbb{R}^d} \|X\beta - Y\|^2$. Then*

$$\|Y\|^2 = \|X\beta^*\|^2 + \|X\beta^* - Y\|^2.$$

Proof. The proof is left as an exercise. □

Sometimes figures help to illustrate a formalism. This is completely unrelated to Theorem 2.

4 Conclusion

Even after centuries of research in the field of data science, there is nothing more versatile than the useful theorem of chapter 3. It is used everywhere and has led to the greatest and most intriguing results, cf. [2]. By the way, the book for the seminar [1] is a great reference and should be cited. Further literature can be found in the respective *Bibliographic Remarks* sections and of course you are welcome to search and add your own references.

Note: BibTeX entries can often be found in the DBLP collection. Google Scholar also offers BibTeX entries, which can be copied into the .bib file and may need some minor adjustments.

References

- [1] S. Shalev-Shwartz and S. Ben-David. *Understanding Machine Learning - From Theory to Algorithms*. Cambridge University Press, 2014.
- [2] J. Someone and J. Someoneelse. Useful theorems, 2003.